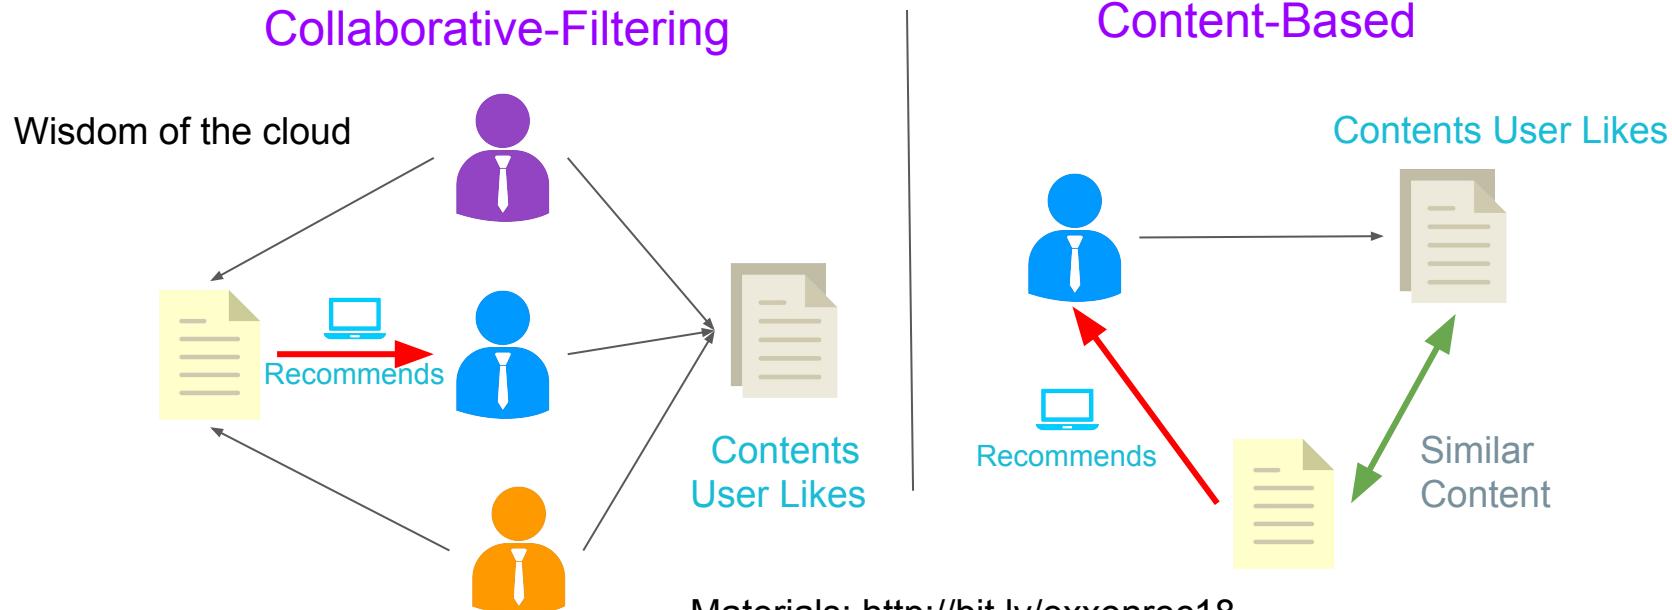


Recommendation System: an overview



Agenda

Overview

Content-based recommendation system

Collaborative filtering systems

Using unstructured information

Metrics

Lab

What is the Goal of a Recommendation System?

- Predict the user's preference toward an item, and recommends it!
- Example: Netflix, Amazon



What do We Know About Our Users?

Explicit Data

- What the Users told us
- Example: Item feedbacks



Sapporo Ichiban Chow Mein, 3.6-Ounce Packages...
Sapporo



I love it

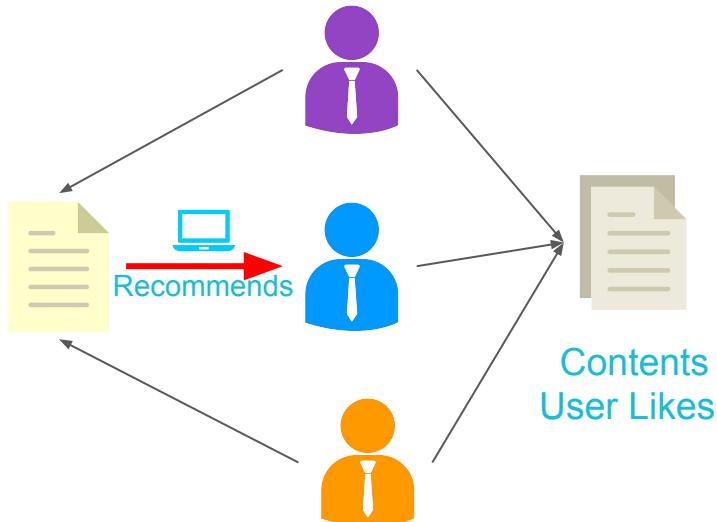
Implicit Data

- What we infer from the user's activity
- Example: Time spent browsing on a page, number of times browsing a page

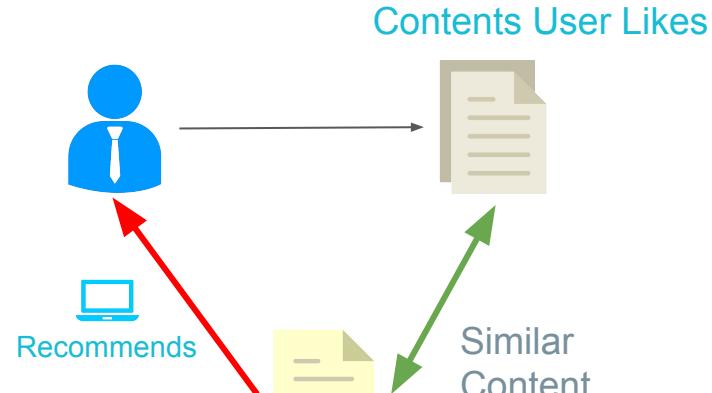


Our Options: Collaborative vs Content-based

Collaborative-Filtering



Content-Based



Icons made by Freepik from www.flaticon.com

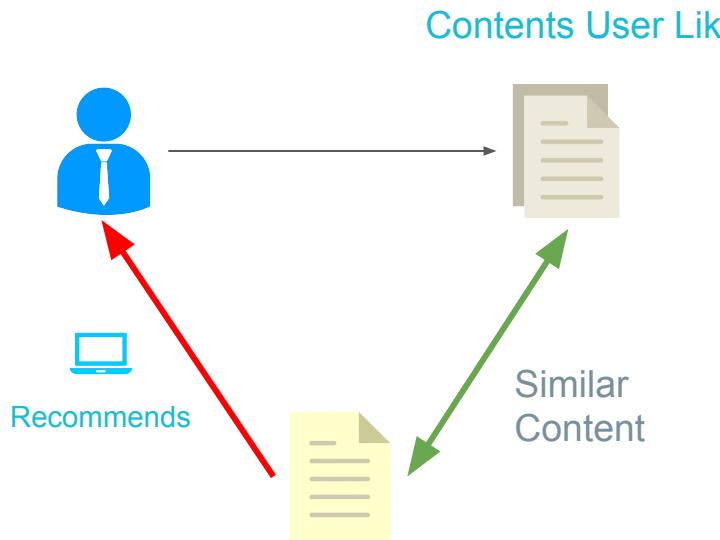
Content-based Recommendation System

- Find an item similar to what the user likes
- Example: Netflix, Amazon

Related to items you've viewed [See more](#)

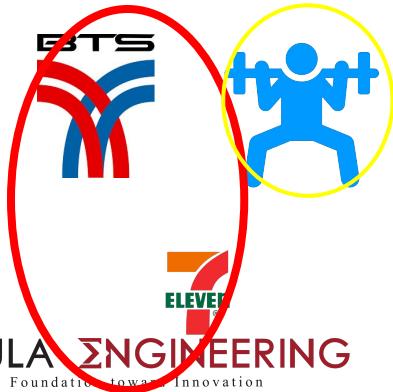
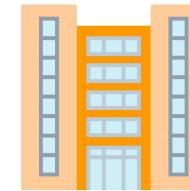
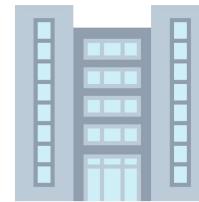
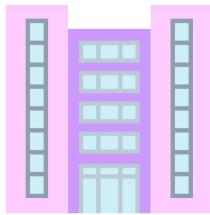


What We Need To Do:

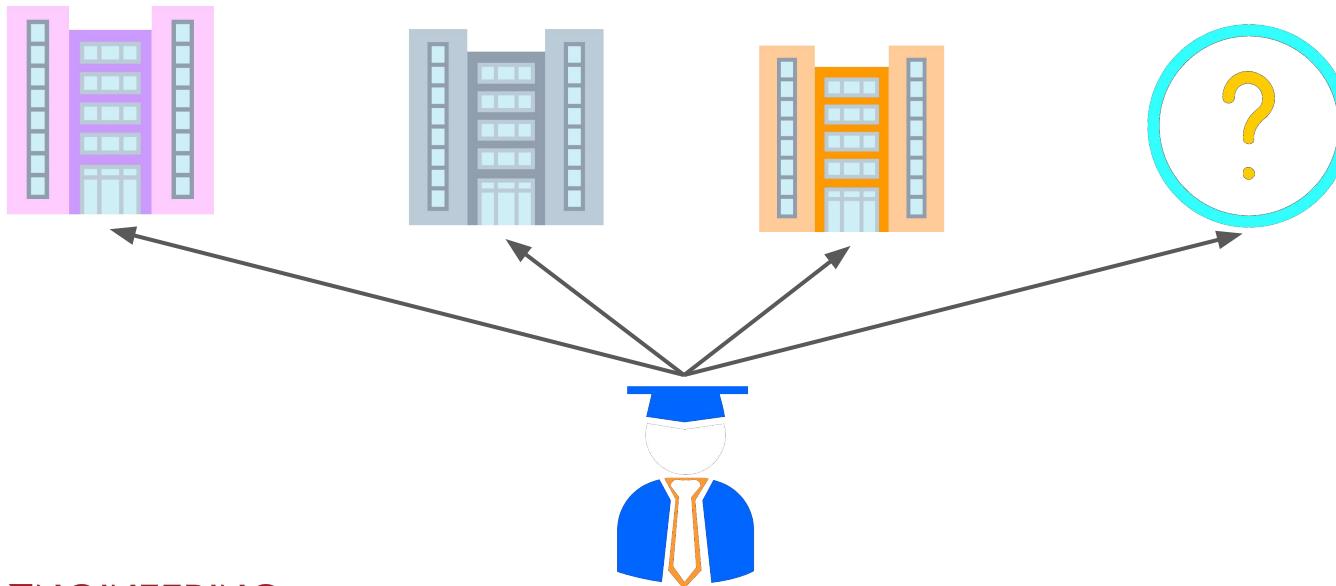


1. Find out what the user likes.
 - Which part of data do we want to use?
 - Rating? Browsing time? Browse count?
2. Construct a User Profile:
 - Which kinds of items does he like?
 - What is the similarity between the items the user likes
3. Match a new item to the User Profile
 - How much will the user like it?
4. Make Prediction!

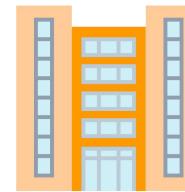
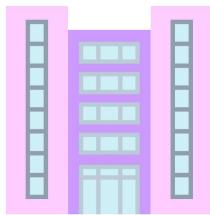
Eyeballing: What Do the Condos Have in Common?



An Intuitive Example: Home browsing

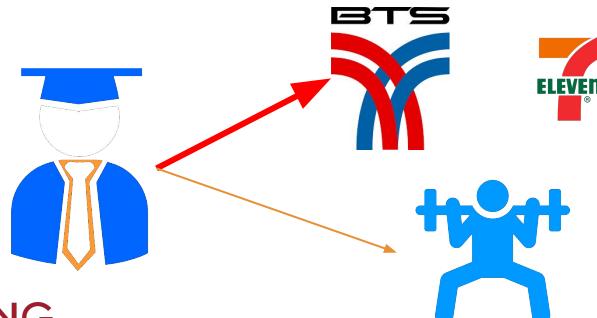


Eyeballing: What Do the Condos Have in Common?

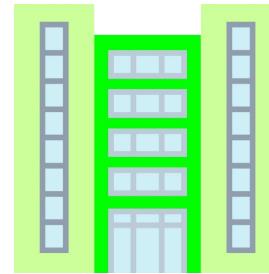
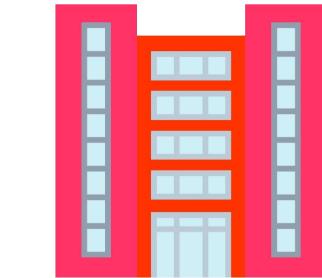
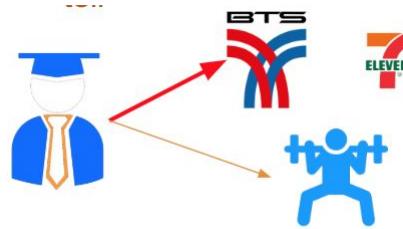


Thus, our customer's profile is:

- Need to be close to the train system
- Need to be close to a convenience store
- Maybe prefer places with fitness center as well, we have too little data to tell

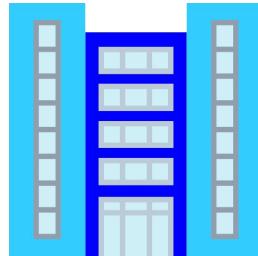
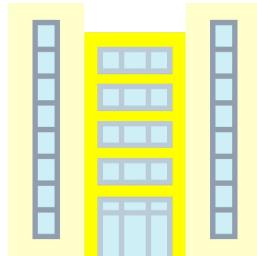
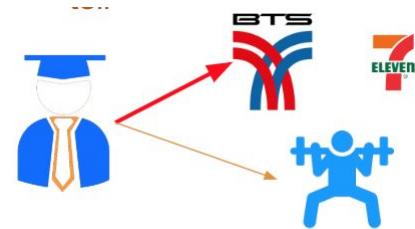


Thus, between



Which should we recommend?

What about?



The answer is not as obvious..

Thus we need a systematic way to tell how well a place fits our model

Item Profile

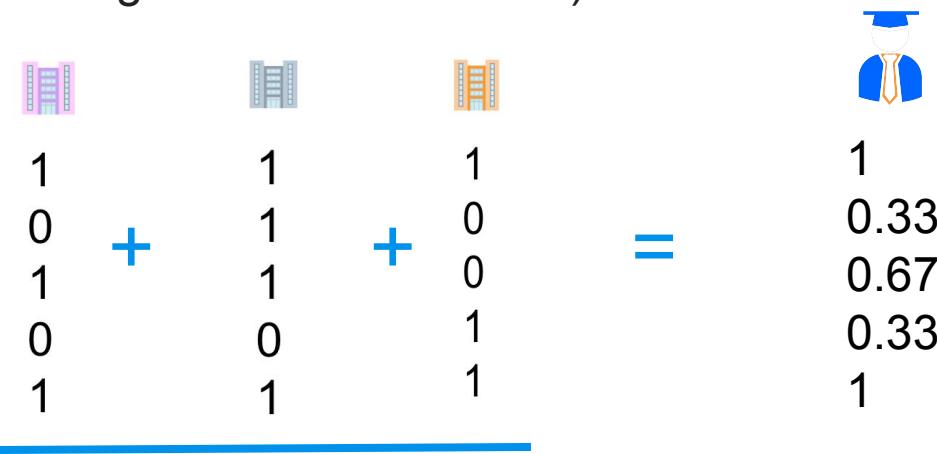
1 = have
0 = don't have



	1	1	1	1	0
A simple icon of a green tree with a brown trunk.	0	1	0	1	0
A blue icon of a person in a squatting position, holding a barbell across their shoulders.	1	1	0	1	0
A blue icon of a swimming pool with white wavy lines representing water.	0	0	1	0	1
The 7-Eleven logo, featuring a red '7' above the word 'ELEVEN' in a white, outlined font.	1	1	1	0	1

Simple Way: we know that the customer have browsed the first three home...

So we can create a user profile by a simple average of the item profile the user browse (Or we can also do weighted average or any more complexed algorithm if we want to!)



Now we can compare which is more similar!

			
	1	1	0
	0.33	1	0
	0.67	1	0
	0.33	0	1
	1	0	1

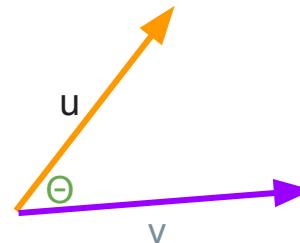
It's still not obvious... but at least we have numbers to work with!

Cosine Similarity

- Our profiles could be viewed as vectors
- We can use the Cosine Similarity to normalize and compare how similar the vectors are

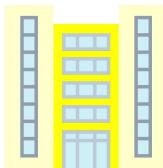
Recall that $u \cdot v = |u||v|\cos\theta$

$$\cos\theta = \frac{u \cdot v}{|u||v|}$$

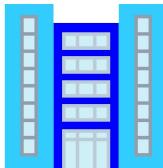


The value of $\cos\theta$ is called the Cosine Similarity, with value closer to one Indicating that the vectors point roughly in the same direction.

Example: Cosine Similarity



$$\cos\theta = \frac{1 \cdot 1 + 0.33 \cdot 1 + 0.67 \cdot 1}{|1.633||\sqrt{3}|} = 0.707$$



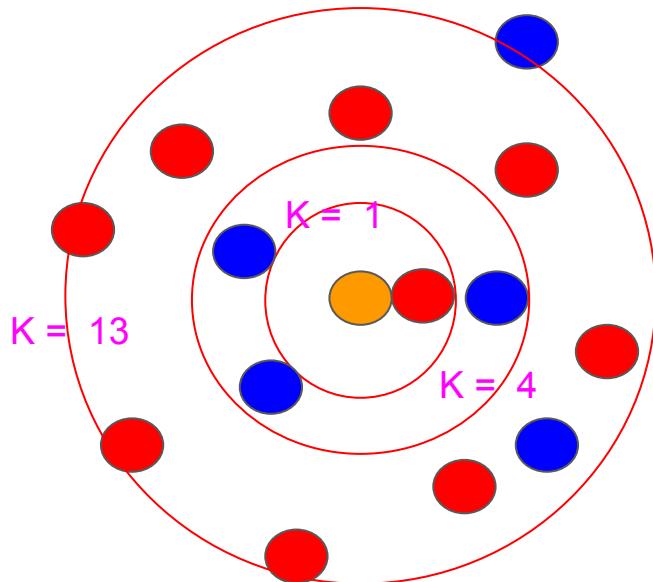
$$\cos\theta = \frac{0.33 \cdot 1 + 1 \cdot 1}{|1.633||\sqrt{3}|} = 0.470$$

The yellow Condo is more similar to what our user has browsed!

Other distance/similarity functions: Euclidean, Jaccard, Earth Mover, etc.

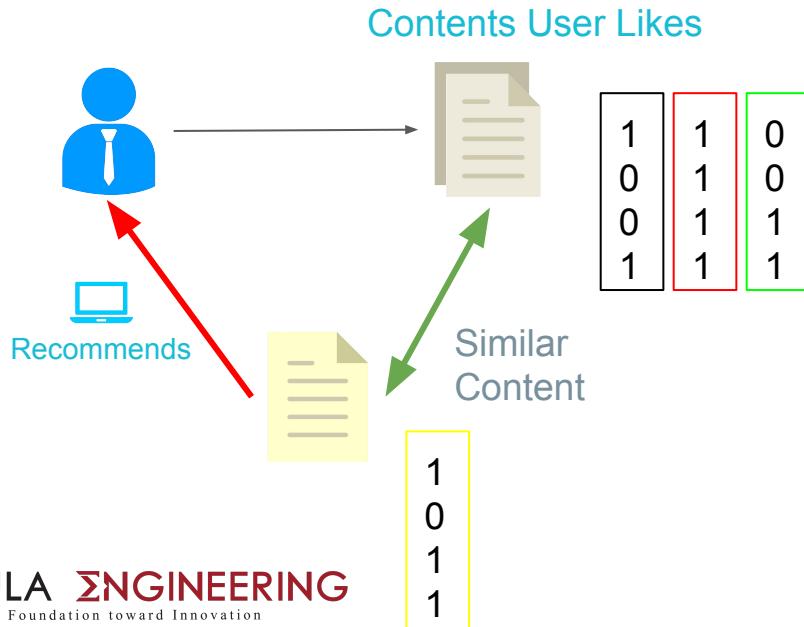
K-Nearest Neighbor

- Classify something based on k items that are most similar to it.



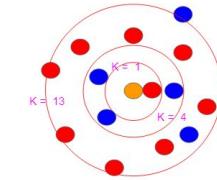
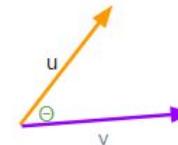
Let's review: Content-based

Content-Based



- Represent items as **features**
- Create a **user profile** based on the items the user interacts with
- **Find** items based on the user profile

	Graduation Cap	Books	Building
1	1	1	0
0.33	0.33	1	0
0.67	0.67	1	0
0.33	0.33	0	1
1	1	0	1



Context information

There's many information available

Ratings/views

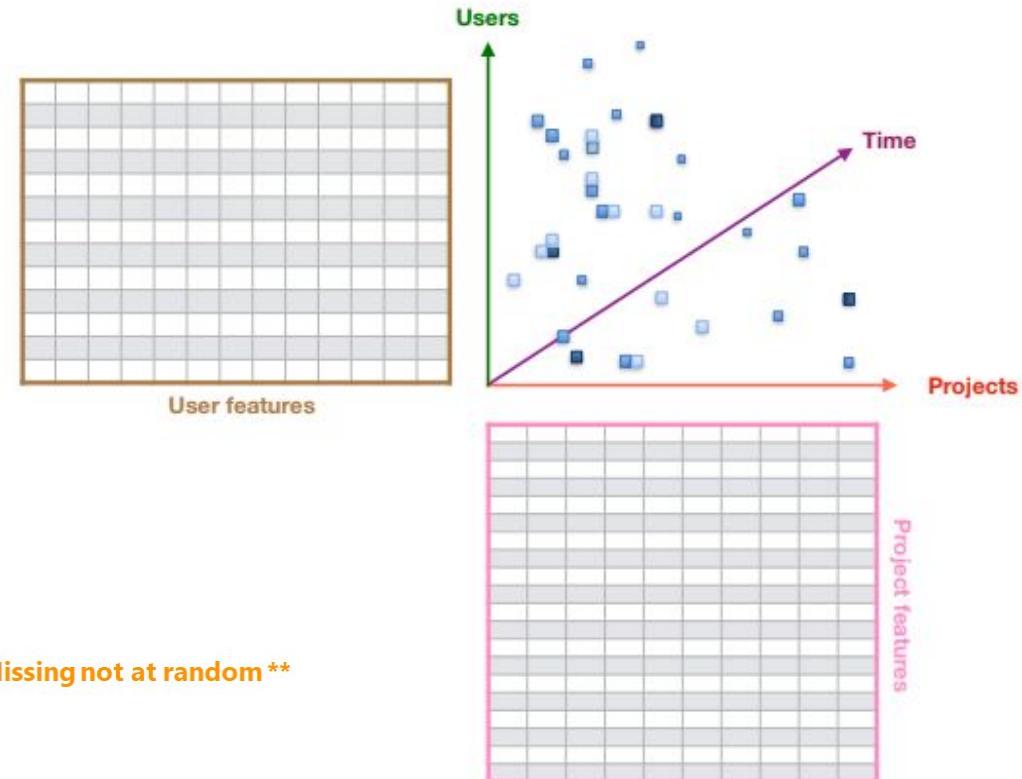
Time

Missing ratings

Side information

User information

Item information



Matrix view (ignoring time)

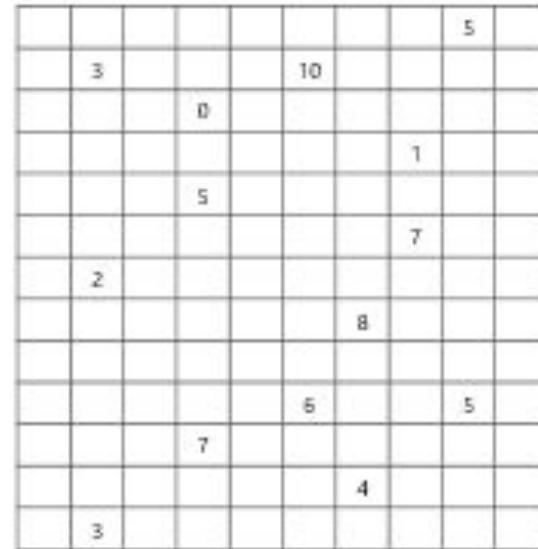
Conventional collaborative filtering tries to fill in the blanks of the view/rating matrices

0/1 matrix

Rating matrix

Implicit info matrix

How to fill?

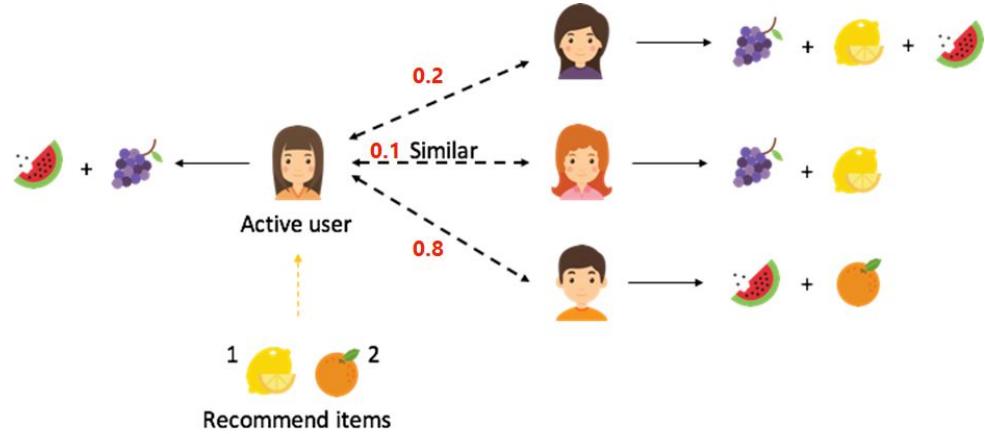


Collaborative filtering methods

- Neighbourhood
- Regression
- Matrix factorization
- Ensemble / Hybrid

Neighborhood models

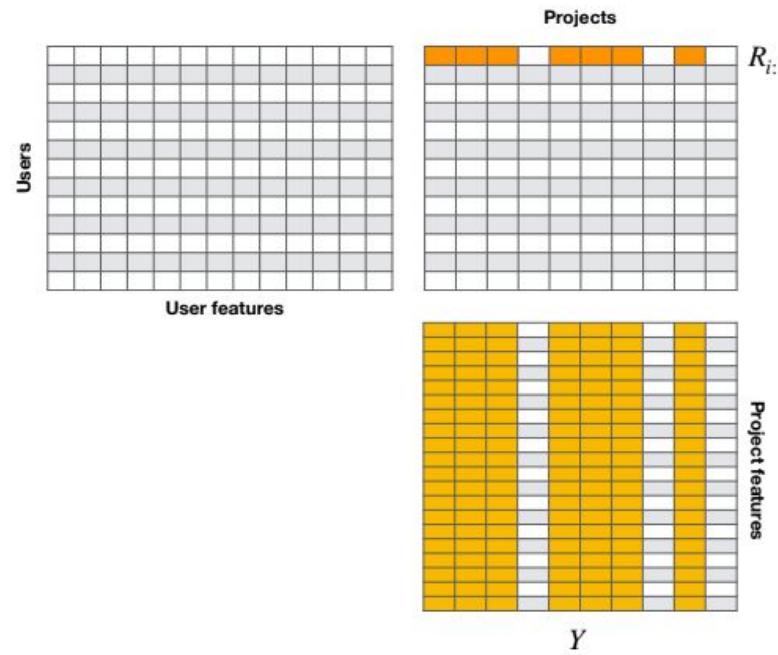
- User-user vs Item-item approach.
- Similarity methods.
- Neighbours aren't independent.



Regression model

1. User-centric

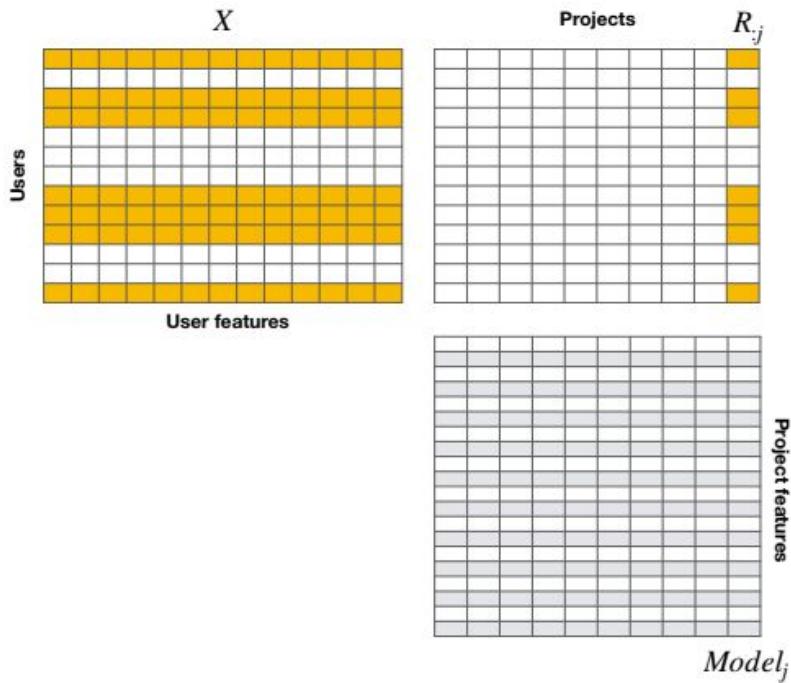
- a. One model per user
- b. Usually ignore missing values (create imbalanced classes)



Regression model

1. User-centric
 - a. One model per user
 - b. Usually ignore missing values (create imbalanced classes)
2. Item-centric
 - a. One model per item
 - b. Same as user-centric

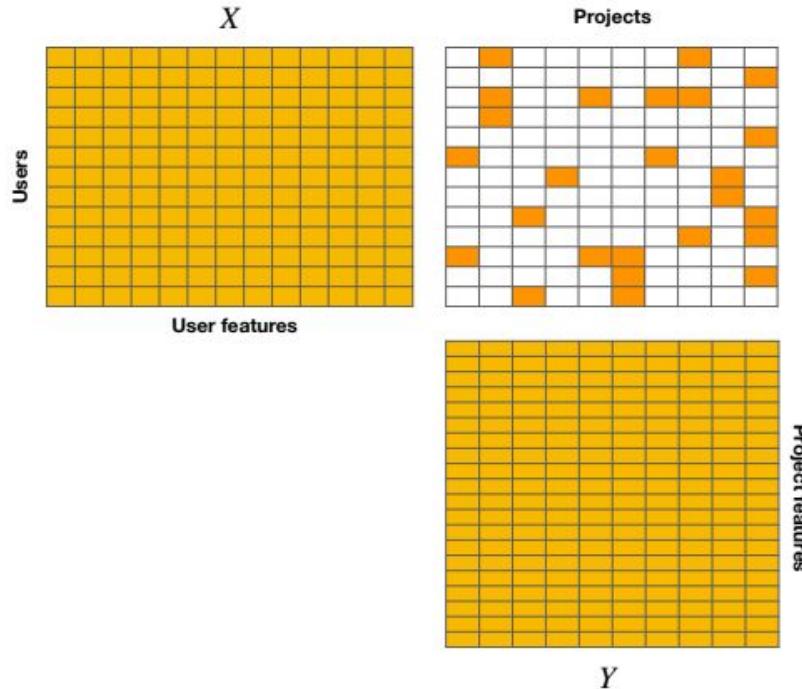
Problems?



Regression model

1. User-centric
 - a. One model per user
 - b. Usually ignore missing values (create imbalanced classes)
2. Item-centric
 - a. One model per item
 - b. Same as user-centric
3. User-item

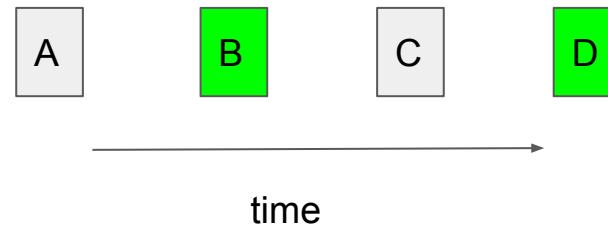
Problems?



Autoregressive model

Modeling time information (sequence)

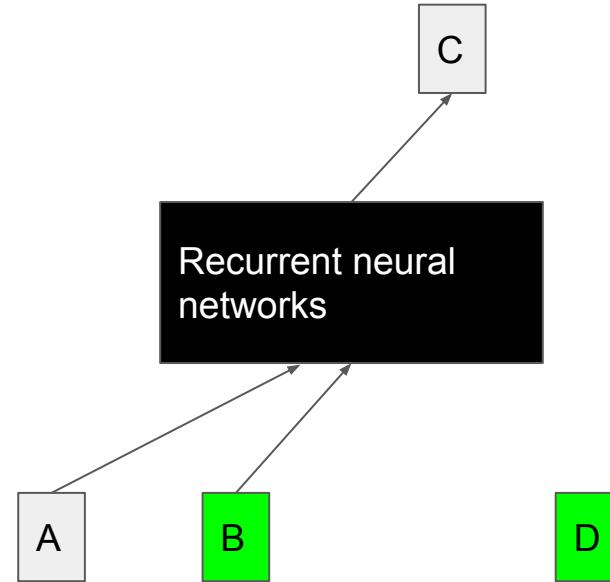
Recurrent Neural Networks



Autoregressive model

Modeling time information (sequence)

Recurrent Neural Networks

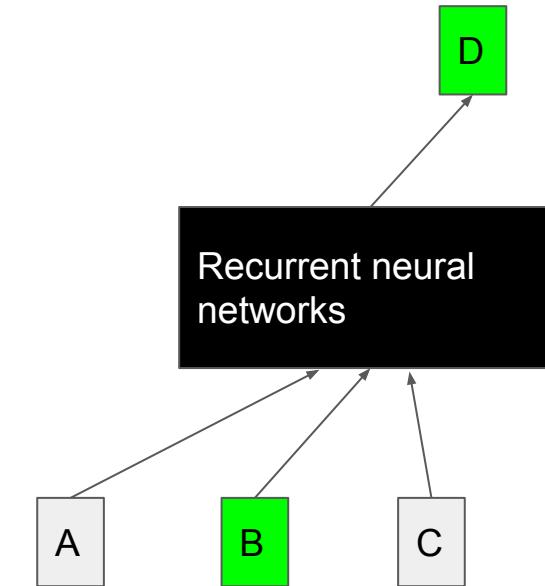


Autoregressive model

Modeling time information (sequence)

Recurrent Neural Networks

Problems?

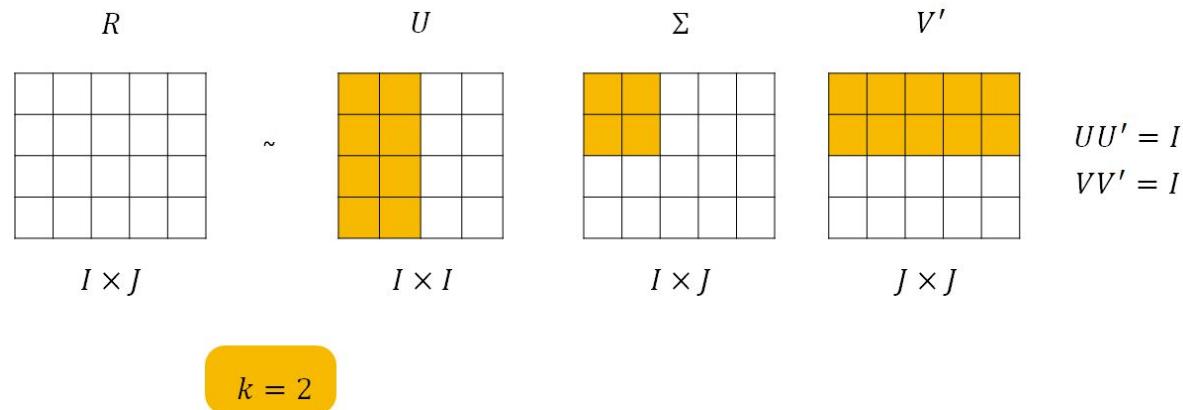


Deep Neural Networks for YouTube Recommendations

Latent Cross: Making Use of Context in Recurrent Recommender Systems

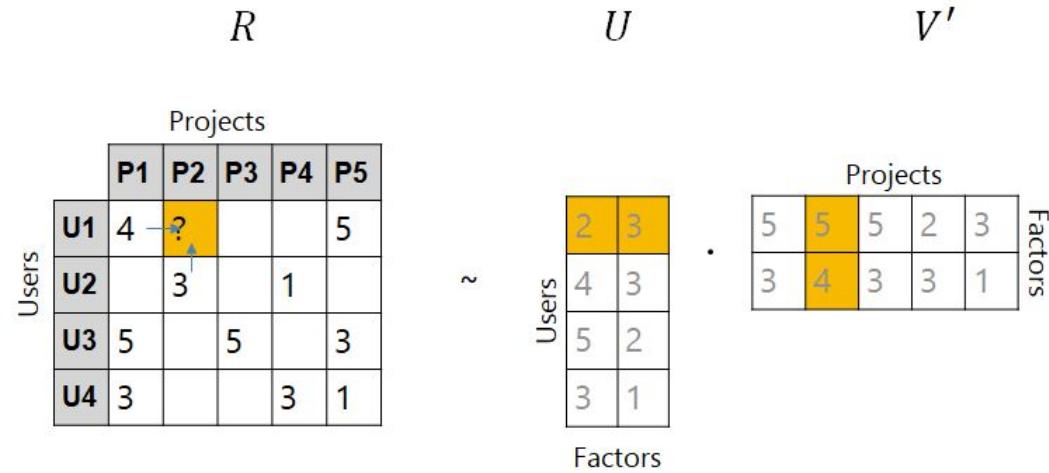
Matrix factorization

Singular Value Decomposition (SVD)



Matrix factorization

Matrix factorization



Objective function:

$$\text{argmin}_{U,V} \sum_{i,j} (R_{ij} - \sum_k U_{ik} V_{jk})^2 + \text{Regularisation}$$

Python Packages:

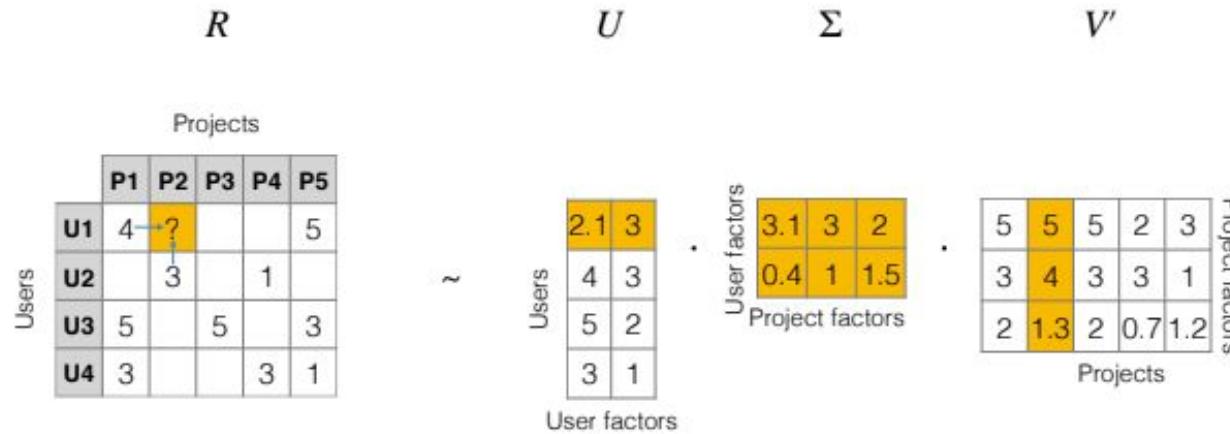
- svds
- nimfa
- Implicit (<https://github.com/benfred/implicit>)
- pyMF (<https://github.com/cthraru/pymf>)

Variants

- Tri-factorisation
- Kernelised Bayesian MF (KBMF):
 - Gönen, Mehmet, Suleiman Khan, and Samuel Kaski. "Kernelized Bayesian matrix factorization." *International Conference on Machine Learning*. 2013.
 - Ammad-Ud-Din, Muhammad, et al. "Integrative and personalized QSAR analysis in cancer by kernelized Bayesian matrix factorization." *Journal of chemical information and modeling* 54.8 (2014): 2347-2359.
- NN meets MF:
 - Dong, Xin, et al. "A Hybrid Collaborative Filtering Model with Deep Structure for Recommender Systems." *AAAI*. 2017.

Tri-factorization

Users and project factors do not need to lie in the same space



Ensemble methods

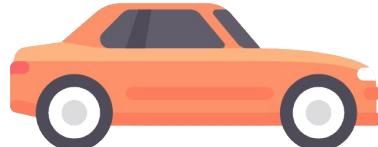
- No single model that use all information
 - Some model works well on certain kind of users
- Weighted combination (linear) or non-linear combination (classifier on top of classifiers)
- Can also train models on errors from previous models

What do we know about the items we want to recommend?

Some items come with easily recognizable attributes

Example: Cars

- Color
- Model
- Seats/Size
- Price
- Eco-friendly?



Others Don't
Example:

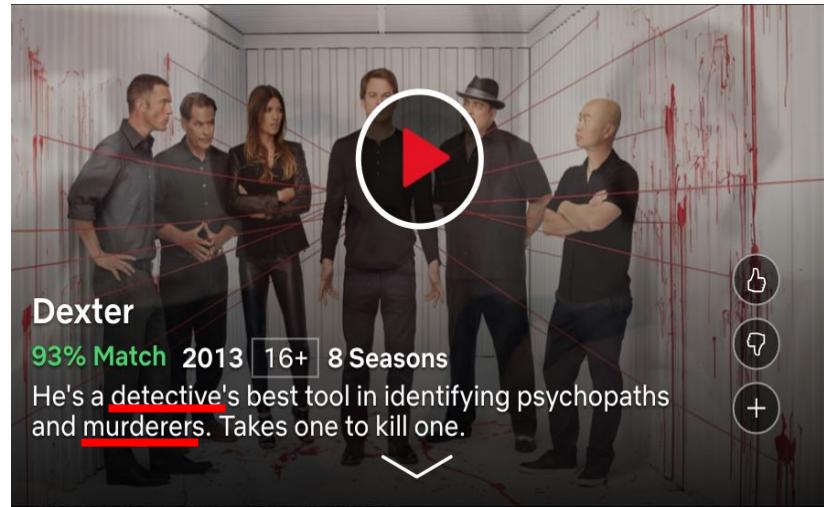
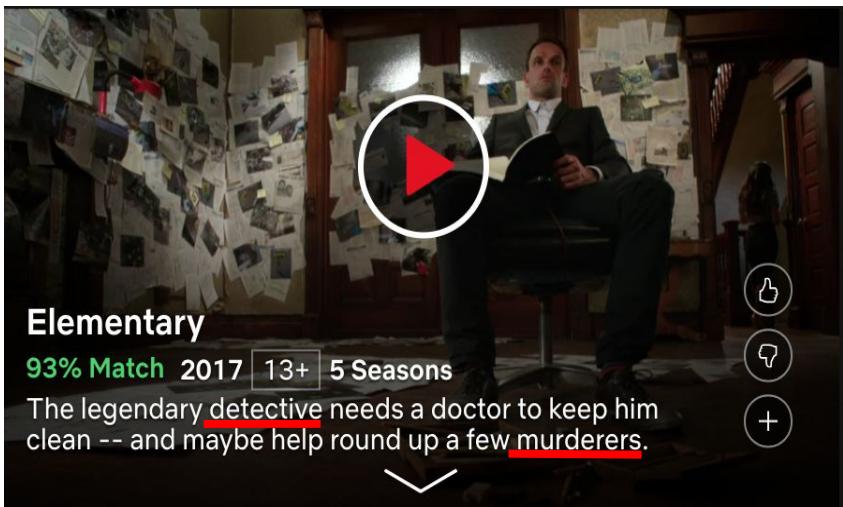
- Wikipedia Pages
- Online Articles



We have both!

Item Description

We can sometimes find related features from the description! (or sometimes name of the item)
Detective+Murderers => likely to be Crime shows

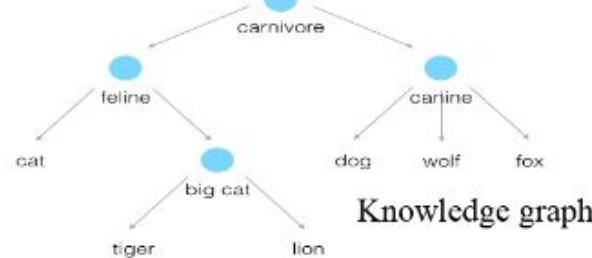


Representing text

Representing text is the most basic task of NLP

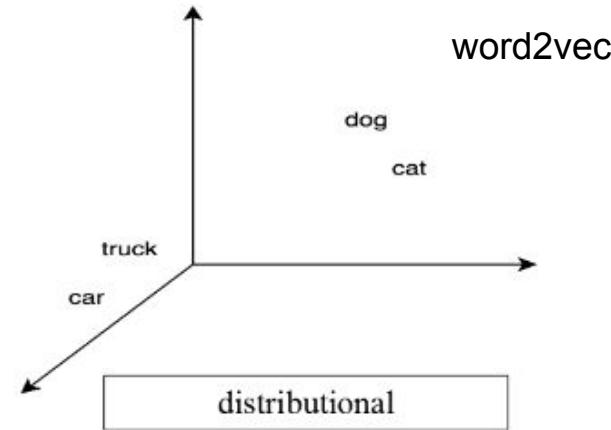
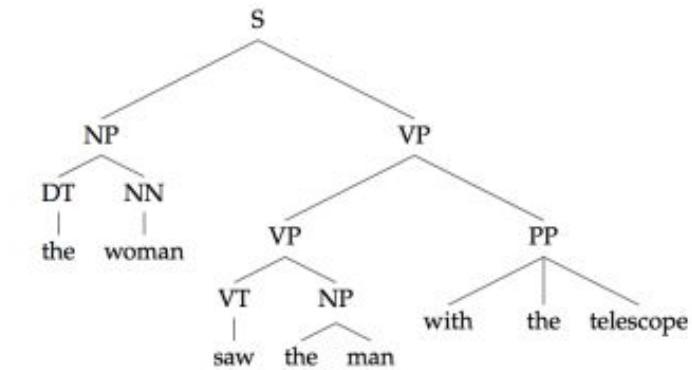
Word level

Sentence/Document level



ເສື່ອ = [0 1 0 0 0 ... 0 0]
One-hot model

symbolic



TF-IDF:

Term Frequency

Inverse Document Frequency

Term Frequency (TF)

= How often does the word
pop up in the description?

Inverse Document Frequency
(IDF)

= How rare is it for a document
to contain that word?

If a term 1) appears a lot in one document and 2)
does not appear a lot in other document, then we
assume it is **significant**.

TF-IDF calculation

- Term Frequency (TF) – per each document

$$TF(w) = \frac{\text{Frequency of word } w \text{ in a document}}{\text{Total number of words in the document}}$$

- Inverse Document Frequency (IDF) – per corpus (all documents)

$$IDF(w) = \log_e\left(\frac{\text{Total number of documents}}{\text{Number of documents that contain word } w}\right)$$

- TF-IDF = TF * IDF

TF-IDF

Word frequency table	อโจเจ้า	มะระ	ໄວອ້ອນ	ກິນ
ບຸພເພສັນນິວາສ	1,000,000			5001
ໄວອ້ອນເຊີບ		3	5000	9999
ໄວອ້ອນແມນ			10000	3781
ໜົກແດງ		100		1090

If a term 1) appears a lot in one document and 2) does not appear a lot in other document, then we assume it is **significant**.

TF-IDF

Word frequency table	อโจเจ้า	มะระ	ໄວອ້ອນ	ກິນ
ບຸພເພສັນນິວາສ	0.01 * 1.4			0.0005 * 0
ໄວອ້ອນເຊີບ		0.02 * 0.7	0.1* 0.7	0.08 * 0
ໄວອ້ອນແມນ			0.1* 0.7	0.04 * 0
ໜົກແດງ		0.06 * 0.7		0.08 * 0

$$TF(w) = \frac{\text{Frequency of word } w \text{ in a document}}{\text{Total number of words in the document}}$$

$$IDF(w) = \log_e\left(\frac{\text{Total number of documents}}{\text{Number of documents that contain word } w}\right)$$

TF-IDF

- We can use TF-IDF to transform item descriptions or title as a feature
- However, consider if there are 100,000 words that we use as features. Would this make our recommender system slow? (Sparse and high dimensional features)
- If we know หรุ and มีระดับ often is included in luxury condos descriptions, maybe we can group descriptions by topics?

Topic modeling

- A description can contain several topic

ค่อนโดยรูส์ไดล์อังกฤษ แห่งแรกในเข้าใหญ่ ที่ติด ณ.ธนารัชต์ มากที่สุด 1 ห้องนอน 1 ห้องน้ำ 1 ห้องนั่งเล่น พร้อมห้องครัวแยกเป็นสัดส่วน

- Convert the proportion of topics into a set of numbers

[0, 0.2, 0, 0.45, 0, 0.35]

Eastern style topic, Western style topic, Facility topic, Layout topic, View topic, Location topic

Supervised topics modeling

If we want to create the representation this way, we need to specify the types of topics and the topics the words in the sentence belongs to.

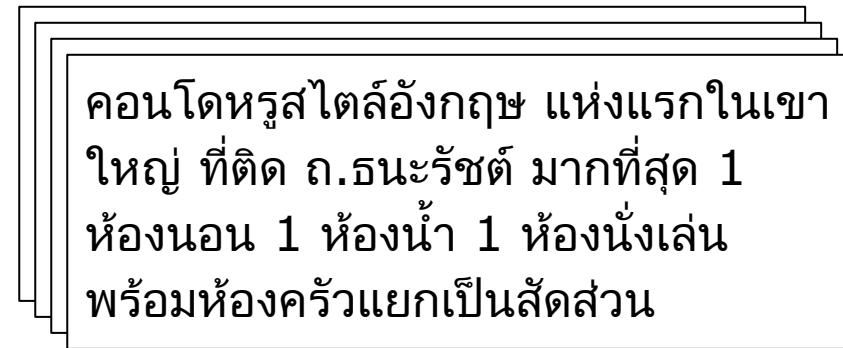
คอนโดหรูสีขาวอังกฤษ แห่งแรกในเข้าใหญ่ ที่ติด ถ.ธนารักษ์ มาก
ที่สุด 1 ห้องนอน 1 ห้องน้ำ 1 ห้องนั่งเล่น พร้อมห้องครัวแยกเป็น^{สัดส่วน}

[0, 0.2, 0, 0.45, 0, 0.35]

Eastern style topic, Western style topic, Facility topic, Layout topic, View topic, Location topic

Unsupervised topic modeling

Can we let the machine learn the topics and the topic assignment by themselves?



Just give it a bunch of descriptions

Key observation: A document/description usually have sparse and coherent topics proportions

Latent Dirichlet Allocation (LDA)

Topics

gene 0.04
dna 0.02
genetic 0.01
...

life 0.02
evolve 0.01
organism 0.01
...

brain 0.04
neuron 0.02
nerve 0.01
...

data 0.02
number 0.02
computer 0.01
...

Documents

Seeking Life's Bare (Genetic) Necessities

COLD SPRING HARBOR, NEW YORK—How many genes does an organism need to survive? Last week at the genome meeting here,* two genome researchers with radically different approaches presented complementary views of the basic genes needed for life. One research team, using computer analyses to compare known genomes, concluded that today's organisms can be sustained with just 250 genes, and that the earliest life forms required a mere 128 genes. The other researcher mapped genes in a simple parasite and estimated that for this organism, 800 genes are plenty to do the job—but that anything short of 100 wouldn't be enough.

Although the numbers don't match precisely, those predictions

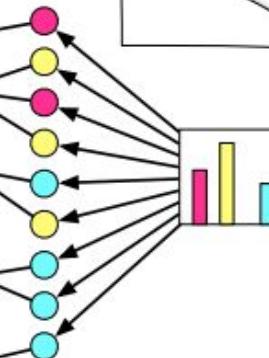
"are not all that far apart," especially in comparison to the 75,000 genes in the human genome, notes Siv Andersson of Uppsala University in Sweden. "I arrived at the 800 number. But coming up with a consensus answer may be more than just a genetic numbers game, particularly as more and more genomes are completely mapped and sequenced." It may be a way of organizing any newly sequenced genome," explains Arcady Mushegian, a computational molecular biologist at the National Center for Biotechnology Information (NCBI) in Bethesda, Maryland. Comparing all



Stripping down. Computer analysis yields an estimate of the minimum modern and ancient genomes.

SCIENCE • VOL. 272 • 24 MAY 1996

Topic proportions and assignments



Automatically
Discovers topics
(Need to specify
the number of
topics)

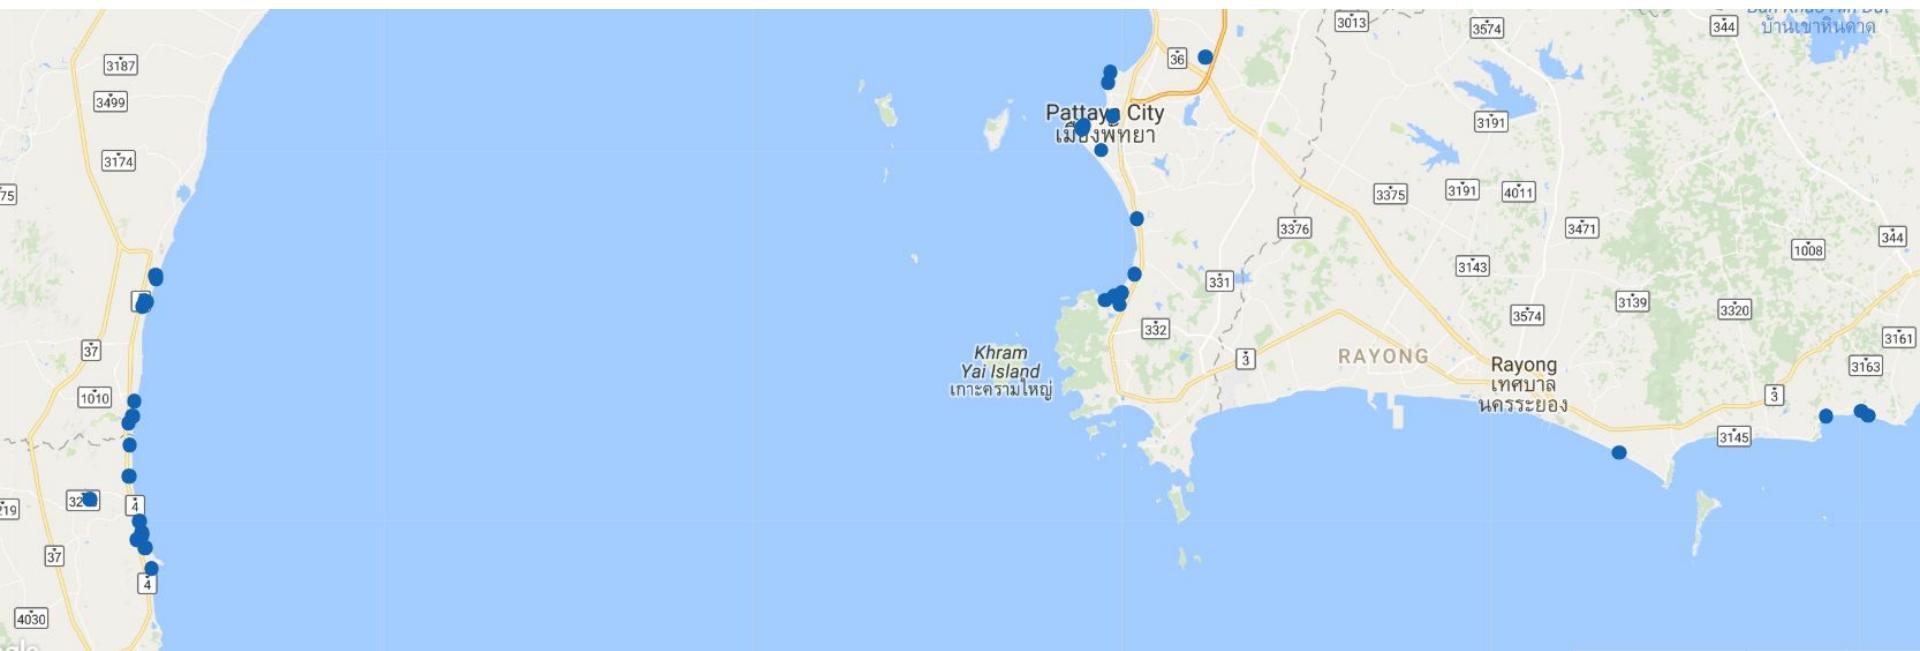
Introduction to Probabilistic Topic Models, Blei 2011

<http://menome.com/wp/wp-content/uploads/2014/12/Blei2011.pdf>

LDA Examples

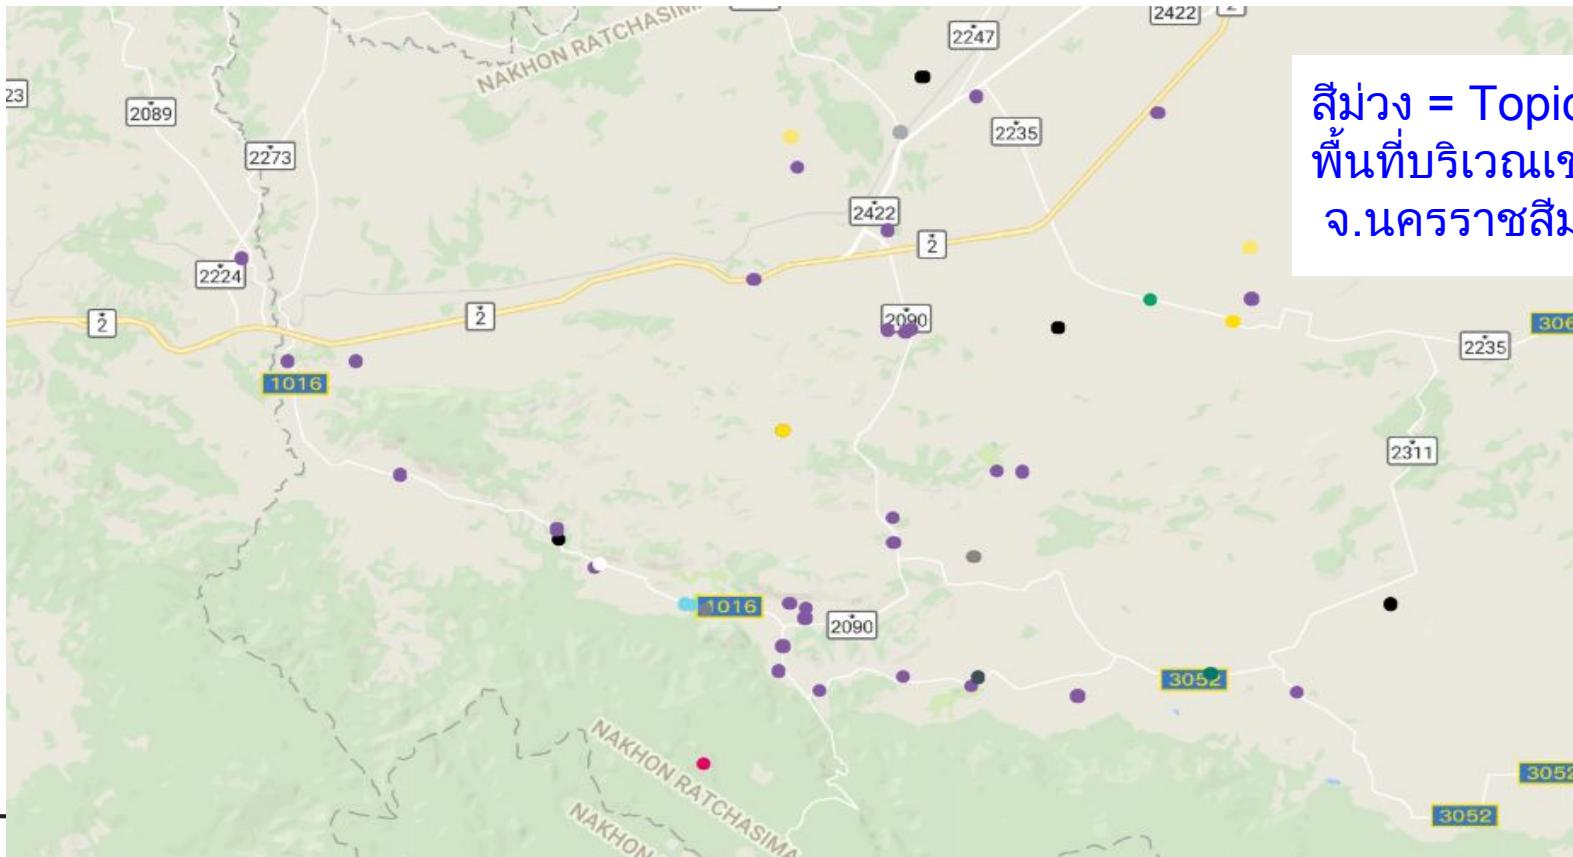
Topic 28

0.068*"วิว" + 0.058*"ทะเล" + 0.038*"คอนโด" + 0.029*"หัว" + 0.027*"คอนโดมิเนียม" + 0.025*"มองเห็น" + 0.023*"ทัศนียภาพ" + 0.022*"ชายหาด"



Topic 9

0.071*"ธรรมชาติ" + 0.031*"บรรยายกาศ" + 0.028*"ร่มรื่น" + 0.027*"บ้าน" + 0.025*"ท่ามกลาง" + 0.025*"สวน" + 0.025*"สัมผัส" + 0.021*"พื้นที่"



ลีม่วง = Topic 9
พื้นที่บริเวณเข้าใหญ่
จ.นครราชสีมา

Topic 40

0.115*"ระดับ" + 0.066*"เห็นอ" + 0.046*"หร" + 0.031*"ทำเล" + 0.026*"ชีวิต" + 0.026*"ใชชีวิต" + 0.016*"สไตล" + 0.016*"สะท้อน"

Topic 17

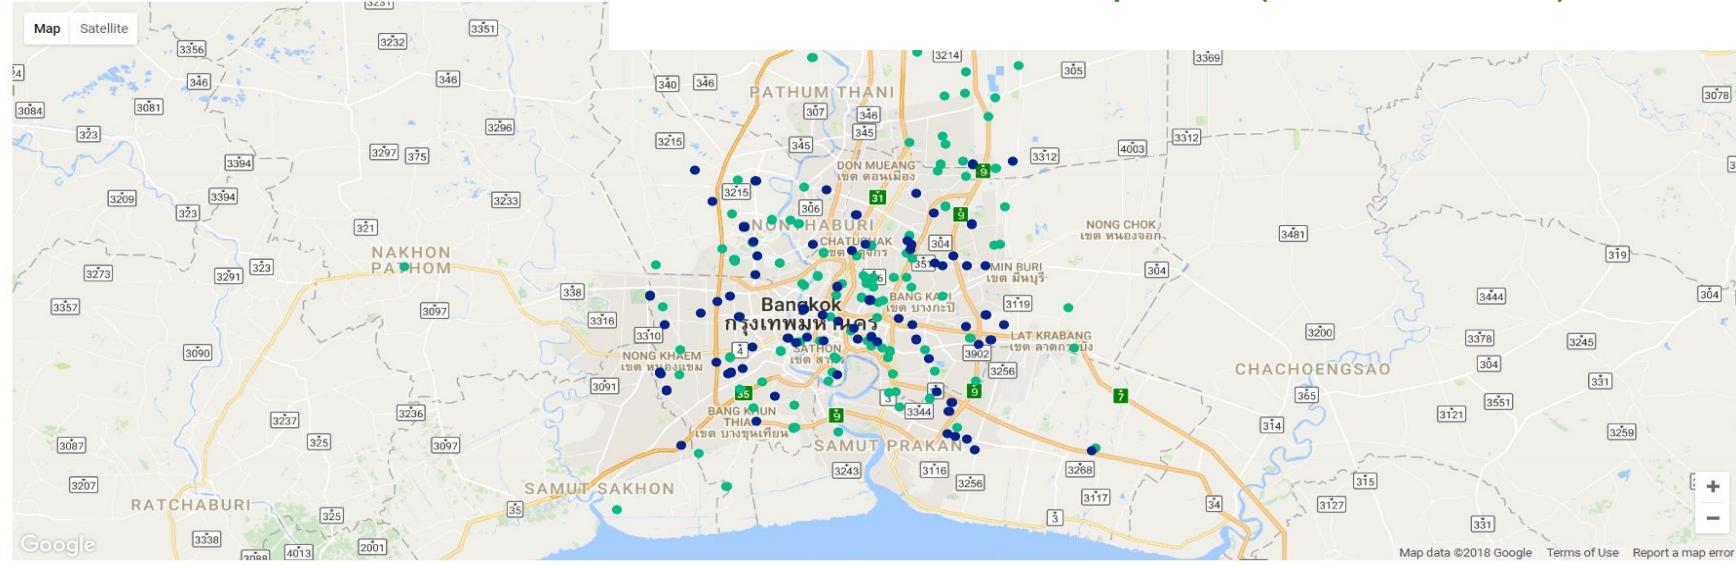
0.077*"พื้นท" + 0.060*"ออกแบบ" + 0.045*"โลง" + 0.039*"โปรดง" + 0.038*"ใชสอย" + 0.020*"ประโยชนช" + 0.018*"หอง" + 0.017*"อาคาร"



Select All Type Unselect All Type default
 บ้านเดี่ยว บ้านแฝด ทาวน์เฮาส คอนโดมิเนียม อาคารพาณิชย โรงแรมไฟฟ ที่ดินเปล่า ทาวน์豪

Select All Cluster Unselect All Cluster
 cluster 0 cluster 1 cluster 2 cluster 3 cluster 4
 cluster 14 cluster 15 cluster 16 cluster 17 cluster 1
 cluster 27 cluster 28 cluster 29 cluster 30 cluster 3

สิน้ำเงินเข้ม = โครงการที่มี Topic 40 อญญา (หร, ระดับ)
สีเขียว = โครงการที่มี Topic 17 (โครงการทัวไป)



Other ways to represent a document

Words to vectors - word2vec

Document to vectors - Averaging word2vec, doc2vec and variants

We do not claim LDA is better than these but it's cool to talk about in the workshop
:)

LDA is interpretable though

Evaluation metrics

Recommendation system commonly used metrics

Best metric is via A/B testing: click through rate, conversion rate, time spent on content, etc.

Mean Square Error (MSE)

Mean Average Precision (MAP@K)

normalized Discounted Cumulative Gain (nDCG)

Evaluation: Precision and Recall

- **Precision** is the **correctness** of our recommendation

Precision = Number of Correct Recommendation / Number Recommended

- **Recall** is the ability of our recommendation to obtain the correct recommendation

Recall = Number of correct answer that get recommended / number of correct answers

Precision and Recall Example

We recommend [H,F,T,**A**,E],

the correct answer is [A,B,C,D]

- We got 1 item(A) correct out of 5 recommended items, so our **precision** is $\frac{1}{5} = 0.2$
- We got 1 item(A) out of 4 correct answers, so our **recall** is $\frac{1}{4} = 0.25$

Evaluation: Average Precision

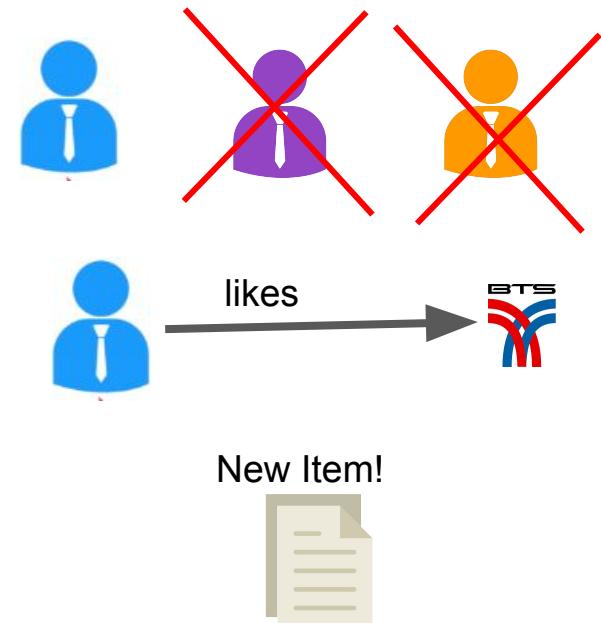
- Average Precision = the Sum of (precision at i) x (the change in recall)
- Example: We recommend [E,K,A,P,O,L], while the correct answer is [A,B,C,D,E]
 - We get two items right at position 1 and 3
 - Precision@1 = 1/1 = 1 Precision@3 = $\frac{2}{3}$ = 0.67
 - Note: Change in recall is always equal to 1/(number of correct items) = $\frac{1}{5}$
 - Average Precision = sum of precision over correct items/number of correct items = $(1+0.67)/5 = 1.67/5 = 0.334$
- What if Instead of [E,K,A,P,O,L] we recommend [E,A,K,P,O,L]
- AP = $(1/1+2/2)/5 = 0.4$ **AP rewards correct rankings!**

Evaluation: Mean Average Precision at K (MAP@K)

- Mean Average Precision is simply the mean of AP for multiple users
- In MAP@K, we only care about the first **K** recommendations
 - You are not penalized for incorrect answers in the K recommendations. Send K answers!
 - Change in recall = $1/\min(K, \text{number correct answer})$
- MAP value both **correctness** of recommendation and **correct rankings**.

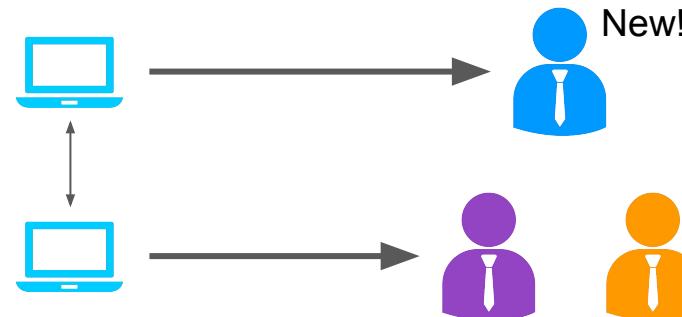
Why should we use Content-based over collaborative filtering?

- We don't need to know about any other users to make a prediction
- Can explain why an item is recommended (What attribute in the item that the user likes)
- Works well on new items with no reviews from other users



Cold Start Problem

- Some recommenders (content-based and collaborative filtering) build user profiles based on past histories
 - We do not have any history for a first time user!
 - Or too little history
- Potential Solution: Hybrid Recommender
 - Some recommenders do not rely on user history(non-personalized recommender etc.)



Cold Start Problem - Item side

- Collaborative filtering methods cannot give suggestions for new item that appears
- Potential Solution: suggestion based on similar items (item features)

Lab overview

1 Data exploration and preprocessing

2 NLP (word segmentation + LDA)

3 Content-based recommendation

4 MF-based recommendation

Hackathon link: <http://bit.ly/homehack18rec>